

Lab Assignment 1: Cache Memories

Neelanjan Goswami

400414867

Part 1: Cache Basics

Solve the following problems and include the solution in the submitted report.

Problem 1

Consider a machine with a byte addressable main memory of 2^8 bytes and block size of 4 bytes. Assume that a direct mapped cache consisting of 8 lines is used with this machine.

1) How is an 8-bit memory address divided into tag, line number, and byte number?

Answer:

To divide an 8-bit memory address into tag, line number, and byte number for a direct-mapped cache with 8 lines and a block size of 4 bytes, you can follow these steps:

Number of Bits for Line Number = $\log_2(\text{Number of Lines})$

Number of Lines = 8 lines

Number of Bits for Line Number = $\log_2(8) = 3$ bits

Number of Bits for Byte Number = $\log_2(\text{Number of Bytes in a Block})$

Number of Bytes in a Block = 4 bytes

Number of Bits for Byte Number = $\log_2(4) = 2$ bits

Number of Bits for Tag = Total Address Bits - (Bits for Line Number + Bits for Byte Number)

Number of Bits for Tag = 8 bits - (3 bits + 2 bits) = 8 bits - 5 bits = 3 bits

Hence,

- 3 bits for the tag

- 3 bits for the line number

- 2 bits for the byte number

2) Into what line would bytes with each of the following addresses be stored?

0001 1011

0011 0100

1101 0000

1010 1010

Answer:

Given the 8-bit memory addresses, we can determine the cache line in which each address would be stored by looking at the line number (index) part of the address. As we determined earlier, the

line number is represented by the 3 bits following the 2 offset bits in the 8-bit address. Here's how you can determine the cache lines for each address:

0001 1011: The line number part is 110, which is 6 in decimal. So, this byte would be stored in cache line 6.

0011 0100: The line number part is 101, which is 5 in decimal. So, this byte would be stored in cache line 5.

1101 0000: The line number part is 100, which is 4 in decimal. So, this byte would be stored in cache line 4.

1010 1010: The line number part is 010, which is 2 in decimal. So, this byte would be stored in cache line 2.

Hence,

0001 1011 – **Line 6**

0011 0100 – **Line 5**

1101 0000 – **Line 4**

1010 1010 – **Line 2**

3) Suppose the byte with address 1010 0001 is stored in the cache. What are the addresses of the other bytes stored along with it?

Answer:

In a cache memory system, a block of data from main memory is stored in a cache line. In this case, the block size is 4 bytes, which means each cache line holds 4 bytes of data.

Given the byte address 1010 0001, we know that this byte is part of a block in main memory. Since the last 2 bits of the address are used as the byte number (offset) within a block, the addresses of the other bytes in the same block will have the same values for the tag and line number parts but will have different values for the byte number part. Here are the addresses of all bytes stored along with 1010 0001 in the same block:

Byte 0: **1010 0000**

Byte 1: **1010 0001** (the original byte)

Byte 2: **1010 0010**

Byte 3: **1010 0011**

So, the addresses of the other bytes stored along with the byte at address 1010 0001 in the cache are 1010 0000, 1010 0010, and 1010 0011.

4) How many total bytes of memory can be stored in the cache?

Answer:

To calculate the total bytes of memory that can be stored in the cache, you can use the following formula:

Total Cache Size = Number of Lines \times Block Size

Total Cache Size = 8 lines \times 4 blocks

Total Cache Size = **32 bytes**

5) Why is the tag also stored in the cache?

Answer:

The tag is stored in the cache to uniquely identify each block of data. It helps in determining whether the desired data is in the cache (cache hit) or not (cache miss). Without the tag, we wouldn't be able to distinguish between different blocks that map to the same cache line.

Problem 2

Consider the following code:

```
cout << "Hello World";  
cin >> a;  
for(i = 0; i < 50; i++)  
cout<<i;
```

1) Give one example of the spatial locality in the code.

Answer:

Spatial locality refers to the concept that when a particular memory location is accessed, the locations near it are likely to be accessed in the near future. In the context of your code, an example of spatial locality is the loop:

```
for(i = 0; i < 50; i++)  
    cout<<i;
```

Here, the variable `i` is being accessed sequentially from `0` to `49`. This is an example of spatial locality because consecutive memory locations (for each value of `i`) are being accessed in a sequence. This is beneficial for performance because modern computer systems anticipate such patterns and pre-load these nearby memory locations into the cache for faster access.

2) Give one example of the temporal locality in the code.

Answer:

Temporal locality refers to the concept that if a particular memory location is accessed, then it's likely to be accessed again in the near future. In the context of your code, an example of temporal locality is the repeated use of the variable `i` in the loop:

```
for(i = 0; i < 50; i++)  
    cout<<i;
```

Here, the variable `i` is being accessed multiple times - once for the comparison `i < 50`, and once for the increment `i++`, and once for printing `cout<<i`. This is an example of temporal locality because the same memory location (the one where `i` is stored) is being accessed repeatedly in a short period of time. This is beneficial for performance because modern computer systems anticipate such patterns and keep recently accessed memory locations in the cache for faster access.

Part 2: Writing the simulator

C++ code file (cachesim_final.cpp), memory_trace.txt and output.txt attached.

Command to compile the C++ file:

```
g++ < filename.cpp > -o < filename >
```

Command to run the C++ file:

```
./< filename> -b [ block size ] -c [ cache size ] -w [ write policy ] -a [ associativity ]
```

(Use 'stride.cpp' for simulation for all of the problems in Part 5 and 6.)

Example commands to run:

```
g++ cachesim_final.cpp -o cachesim_final
```

```
./cachesim_final -b 32 -c 16384 -w b -a 4 < memory_trace.txt
```

Simulation output of the example:

```

neelanjangoswami@Neelanjans-Air Downloads % ./cachesim_final -b 32 -c 16384 -w b -a 4 < memory_trace.txt
16KB 4-way associative cache:
  Block size = 32 bytes
  Number of [sets,blocks] = [128,512]
  Extra space for tag storage = 768 bytes (4.69%)
  Bits for [tag,index,offset] = [12, 7, 5] = 24
  Write policy = Write-through

```

Hex	Binary	Address	Set	Blk	Memory				
Address	(tag/index/offset)		Tag	Index	off	Way	UWay	Read	Write
c0e7R	0000001100	00000111 00111	12	7	7	-1	0	1	0
376W	0000000000	00011011 10110	0	27	22	-1	0	0	1
c0efR	0000001100	00000111 01111	12	7	15	0	-2	1	0
37eW	0000000000	00011011 11110	0	27	30	0	0	0	1
c0f7R	0000001100	00000111 10111	12	7	23	0	-2	1	0
386W	0000000000	00011100 00110	0	28	6	-1	0	0	1
c0ffR	0000001100	00000111 11111	12	7	31	0	-2	1	0
38eW	0000000000	00011100 01110	0	28	14	0	0	0	1
c107R	0000001100	00001000 00111	12	8	7	-1	0	1	0
396W	0000000000	00011100 10110	0	28	22	0	0	0	1
c10fR	0000001100	00001000 01111	12	8	15	0	-2	1	0
39eW	0000000000	00011100 11110	0	28	30	0	0	0	1
c117R	0000001100	00001000 10111	12	8	23	0	-2	1	0
3a6W	0000000000	00011101 00110	0	29	6	-1	0	0	1
c11fR	0000001100	00001000 11111	12	8	31	0	-2	1	0
3aeW	0000000000	00011101 01110	0	29	14	0	0	0	1

```

nref=16, nread=11, nwrite=8
hits = 11, hit rate = 0.69
misses = 5, miss rate = 0.31
main memory reads=11, main memory writes=8

```

Part 3: Spatial Locality

For the following logic, where A and B are arrays of integer elements:

sum=0 // sum is int

for(i = 0; i < num_elements; i++)

sum=sum+A[i]+B[i]

1. Assume that A[0] is placed in address 0x0 and B[0] is in address 0x100

2. (Paper and Pen Analysis): Experiment with num_elements=8, a four-way set associative cache of size 256-B and three different cache line sizes (16, 8, and 4B). For each configuration, obtain the hit rate and classify the misses, record your observations.

Answer:

Assuming address bit length= 24 bit

Size of int = 4 bytes

cache size = 256 bytes

num_elements=8 and a four-way set associative cache of size 256 bytes.

Cache line sizes: 16B, 8B, and 4B

Number of sets = Cache size / (Cache line size * Number of ways)

For a four-way set associative cache, we have 4 ways.

Now, let's calculate the hit rate and classify the misses for each configuration:

Case 1: Cache Line Size: 4B

Number of cache lines = Cache size / Cache line size = 256B / 4B = 64 cache lines

No. of sets = 64 / 4 = 16

Bits for [tag , index , offset] = [18 , 4 , 2] = 24

Extra space for tag storage = 144B

Variable	Hex Address [tag, index, offset]			Decimal
	tag	index	offset	
	18	4	2	
A[0]	000000 0000 0000 0000	0000	00	0
A[1]	000000 0000 0000 0000	0001	00	4
A[2]	000000 0000 0000 0000	0010	00	8
A[3]	000000 0000 0000 0000	0011	00	12
A[4]	000000 0000 0000 0000	0100	00	16
A[5]	000000 0000 0000 0000	0101	00	20
A[6]	000000 0000 0000 0000	0110	00	24
A[7]	000000 0000 0000 0000	0111	00	28

Variable	Hex Address [tag, index, offset]			Decimal
	tag	index	offset	
	18	4	2	
B[0]	000000 0000 0000 0100	0000	00	256
B[1]	000000 0000 0000 0100	0001	00	260
B[2]	000000 0000 0000 0100	0010	00	264
B[3]	000000 0000 0000 0100	0011	00	268
B[4]	000000 0000 0000 0100	0100	00	272
B[5]	000000 0000 0000 0100	0101	00	276
B[6]	000000 0000 0000 0100	0110	00	280
B[7]	000000 0000 0000 0100	0111	00	284

All misses here are Compulsory Misses

Variable	Operation: [Variable Address]	Hit/Miss		Variable	Operation: [Variable Address]	Hit/Miss
A[0]	R: 0x0	Miss		B[0]	R: 0x100	Miss
A[1]	R: 0x4	Miss		B[1]	R: 0x104	Miss
A[2]	R: 0x8	Miss		B[2]	R: 0x108	Miss
A[3]	R: 0xC	Miss		B[3]	R: 0x10C	Miss
A[4]	R: 0x10	Miss		B[4]	R: 0x110	Miss
A[5]	R: 0x14	Miss		B[5]	R: 0x114	Miss

A[6]	R: 0x18	Miss		B[6]	R: 0x118	Miss
A[7]	R: 0x1C	Miss		B[7]	R: 0x11C	Miss

Total traces (read count) = 16

Hits = 0

Misses = 16 (Compulsory Misses)

Memory reads = 0

Hit rate = 0 / 16 = 0

Case 2: Cache Line Size: 8B

Number of cache lines = Cache size / Cache line size = 256B / 8B = 32 cache lines

No. of sets = 32 / 4 = 8

Bits for [tag , index , offset] = [18 , 3 , 3] = 24

Extra space for tag storage = 72B

Variable	Hex Address [tag, index, offset]			Decimal
	tag	index	offset	
	18	3	3	
A[0]	000000 0000 0000 0000	000	000	0
A[1]	000000 0000 0000 0000	000	100	4
A[2]	000000 0000 0000 0000	001	000	8
A[3]	000000 0000 0000 0000	001	100	12
A[4]	000000 0000 0000 0000	010	000	16
A[5]	000000 0000 0000 0000	010	100	20
A[6]	000000 0000 0000 0000	011	000	24
A[7]	000000 0000 0000 0000	011	100	28

Variable	Hex Address [tag, index, offset]			Decimal
	tag	index	offset	
	18	3	3	
B[0]	000000 0000 0000 0100	000	000	256
B[1]	000000 0000 0000 0100	000	100	260
B[2]	000000 0000 0000 0100	001	000	264
B[3]	000000 0000 0000 0100	001	100	268
B[4]	000000 0000 0000 0100	010	000	272
B[5]	000000 0000 0000 0100	010	100	276
B[6]	000000 0000 0000 0100	011	000	280
B[7]	000000 0000 0000 0100	011	100	284

All misses here are Compulsory Misses

Variable	Operation: [Variable Address]	Hit/Miss		Variable	Operation: [Variable Address]	Hit/Miss
A[0]	R: 0x0	Miss		B[0]	R: 0x100	Miss
A[1]	R: 0x4	Hit		B[1]	R: 0x104	Hit
A[2]	R: 0x8	Miss		B[2]	R: 0x108	Miss
A[3]	R: 0xC	Hit		B[3]	R: 0x10C	Hit
A[4]	R: 0x10	Miss		B[4]	R: 0x110	Miss
A[5]	R: 0x14	Hit		B[5]	R: 0x114	Hit
A[6]	R: 0x18	Miss		B[6]	R: 0x118	Miss
A[7]	R: 0x1C	Hit		B[7]	R: 0x11C	Hit

Total traces (read count) = 16

Hits = 8

Misses = 8 (Compulsory Misses)

Memory reads = 8

Hit rate = $8 / 16 = 0.5$

Case 3: Cache Line Size: 16B

Number of cache lines = Cache size / Cache line size = 256B / 16B = 16 cache lines

No. of sets = 16 / 4 = 4

Bits for [tag , index , offset] = [18 , 2 , 4] = 24

Extra space for tag storage = 36B

Variable	Hex Address [tag, index, offset]			Decimal
	tag	index	offset	
	18	2	4	
A[0]	000000 0000 0000 0000	00	0000	0
A[1]	000000 0000 0000 0000	00	0100	4
A[2]	000000 0000 0000 0000	00	1000	8
A[3]	000000 0000 0000 0000	00	1100	12
A[4]	000000 0000 0000 0000	01	0000	16
A[5]	000000 0000 0000 0000	01	0100	20
A[6]	000000 0000 0000 0000	01	1000	24
A[7]	000000 0000 0000 0000	01	1100	28

Variable	Hex Address [tag, index, offset]			Decimal
	tag	index	offset	
	18	2	4	
B[0]	000000 0000 0000 0100	00	0000	256
B[1]	000000 0000 0000 0100	00	0100	260

B[2]	000000 0000 0000 0100	00	1000	264
B[3]	000000 0000 0000 0100	00	1100	268
B[4]	000000 0000 0000 0100	01	0000	272
B[5]	000000 0000 0000 0100	01	0100	276
B[6]	000000 0000 0000 0100	01	1000	280
B[7]	000000 0000 0000 0100	01	1100	284

All misses here are Compulsory Misses

Variable	Operation: [Variable Address]	Hit/Miss		Variable	Operation: [Variable Address]	Hit/Miss
A[0]	R: 0x0	Miss		B[0]	R: 0x100	Miss
A[1]	R: 0x4	Hit		B[1]	R: 0x104	Hit
A[2]	R: 0x8	Hit		B[2]	R: 0x108	Hit
A[3]	R: 0xC	Hit		B[3]	R: 0x10C	Hit
A[4]	R: 0x10	Miss		B[4]	R: 0x110	Miss
A[5]	R: 0x14	Hit		B[5]	R: 0x114	Hit
A[6]	R: 0x18	Hit		B[6]	R: 0x118	Hit
A[7]	R: 0x1C	Hit		B[7]	R: 0x11C	Hit

Total traces (read count) = 16

Hits = 12

Misses = 4 (Compulsory Misses)

Memory reads = 4

Hit rate = 12 / 16 = 0.75

Hence, Hit rates for 4 byte cache line = 0, for 8 byte cache line = 0.5 & for 16 byte cache line = 0.75

3. (Simulation): hand craft the traces for each of the configurations in Step 2 and run this through the simulator your created in Part 2.

Answer:

Here are the traces for the simulation:

R:0

R:4

R:8

R:C

R:10

R:14

R:18

R:1C

R:100
R:104
R:108
R:10C
R:110
R:114
R:118
R:11C

Case 1:

```
neelanjangoswami@Neelanjans-Air Downloads % ./cachesim_final -b 4 -c 256 -w r -a 4 < memory_trace.txt
0KB 4-way associative cache:
  Block size = 4 bytes
  Number of [sets,blocks] = [16,64]
  Extra space for tag storage = 144 bytes (56.25%)
  Bits for [tag,index,offset] = [18, 4, 2] = 24
  Write policy = Write-back

Hex  Binary Address      Set  Blk  Memory
Address (tag/index/offset)  Tag Index off  Way UWay Read Write
=====
0R 0000000000 00000000 00000 0 0 0 -1 0 1 0
4R 0000000000 00000001 00000 0 1 0 -1 0 1 0
8R 0000000000 00000010 00000 0 2 0 -1 0 1 0
cR 0000000000 00000011 00000 0 3 0 -1 0 1 0
10R 0000000000 00000100 00000 0 4 0 -1 0 1 0
14R 0000000000 00000101 00000 0 5 0 -1 0 1 0
18R 0000000000 00000110 00000 0 6 0 -1 0 1 0
1cR 0000000000 00000111 00000 0 7 0 -1 0 1 0
100R 0000000100 00000000 00000 4 0 0 -1 0 1 0
104R 0000000100 00000001 00000 4 1 0 -1 0 1 0
108R 0000000100 00000010 00000 4 2 0 -1 0 1 0
10cR 0000000100 00000011 00000 4 3 0 -1 0 1 0
110R 0000000100 00000100 00000 4 4 0 -1 0 1 0
114R 0000000100 00000101 00000 4 5 0 -1 0 1 0
118R 0000000100 00000110 00000 4 6 0 -1 0 1 0
11cR 0000000100 00000111 00000 4 7 0 -1 0 1 0
nref=16, nread=16, nwrite=0
hits = 0, hit rate = 0.00
misses = 16, miss rate = 1.00
main memory reads=16, main memory writes=0
```

Case 2:

```
neelanjangoswami@Neelanjans-Air Downloads % ./cachesim_final -b 8 -c 256 -w r -a 4 < memory_trace.txt
0KB 4-way associative cache:
  Block size = 8 bytes
  Number of [sets,blocks] = [8,32]
  Extra space for tag storage = 72 bytes (28.12%)
  Bits for [tag,index,offset] = [18, 3, 3] = 24
  Write policy = Write-back

Hex  Binary Address      Set  Blk  Memory
Address (tag/index/offset)  Tag Index off  Way UWay Read Write
=====
0R 0000000000 00000000 00000 0 0 0 -1 0 1 0
4R 0000000000 00000000 00100 0 0 4 0 -2 1 0
8R 0000000000 00000001 00000 0 1 0 -1 0 1 0
cR 0000000000 00000001 00100 0 1 4 0 -2 1 0
10R 0000000000 00000010 00000 0 2 0 -1 0 1 0
14R 0000000000 00000010 00100 0 2 4 0 -2 1 0
18R 0000000000 00000011 00000 0 3 0 -1 0 1 0
1cR 0000000000 00000011 00100 0 3 4 0 -2 1 0
100R 0000000100 00000000 00000 4 0 0 -1 0 1 0
104R 0000000100 00000000 00100 4 0 4 0 -2 1 0
108R 0000000100 00000001 00000 4 1 0 -1 0 1 0
10cR 0000000100 00000001 00100 4 1 4 0 -2 1 0
110R 0000000100 00000010 00000 4 2 0 -1 0 1 0
114R 0000000100 00000010 00100 4 2 4 0 -2 1 0
118R 0000000100 00000011 00000 4 3 0 -1 0 1 0
11cR 0000000100 00000011 00100 4 3 4 0 -2 1 0
nref=16, nread=16, nwrite=0
hits = 8, hit rate = 0.50
misses = 8, miss rate = 0.50
main memory reads=16, main memory writes=0
```

Case 3:

```
neelanjangoswami@Neelanjans-Air Downloads % ./cachesim_final -b 16 -c 256 -w r -a 4 < memory_trace.txt
0KB 4-way associative cache:
  Block size = 16 bytes
  Number of [sets,blocks] = [4,16]
  Extra space for tag storage = 36 bytes (14.06%)
  Bits for [tag,index,offset] = [18, 2, 4] = 24
  Write policy = Write-back

Hex Binary Address          Set Blk   Memory
Address (tag/index/offset)  Tag Index off  Way UWay Read Write
=====
0R 0000000000 00000000 00000 0 0 0 -1 0 1 0
4R 0000000000 00000000 00100 0 0 4 0 -2 1 0
8R 0000000000 00000000 01000 0 0 8 0 -2 1 0
cR 0000000000 00000000 01100 0 0 12 0 -2 1 0
10R 0000000000 00000001 00000 0 1 0 -1 0 1 0
14R 0000000000 00000001 00100 0 1 4 0 -2 1 0
18R 0000000000 00000001 01000 0 1 8 0 -2 1 0
1cR 0000000000 00000001 01100 0 1 12 0 -2 1 0
100R 0000000100 00000000 00000 4 0 0 -1 0 1 0
104R 0000000100 00000000 00100 4 0 4 0 -2 1 0
108R 0000000100 00000000 01000 4 0 8 0 -2 1 0
10cR 0000000100 00000000 01100 4 0 12 0 -2 1 0
110R 0000000100 00000001 00000 4 1 0 -1 0 1 0
114R 0000000100 00000001 00100 4 1 4 0 -2 1 0
118R 0000000100 00000001 01000 4 1 8 0 -2 1 0
11cR 0000000100 00000001 01100 4 1 12 0 -2 1 0
nref=16, nread=16, nwrite=0
hits = 12, hit rate = 0.75
misses = 4, miss rate = 0.25
main memory reads=16, main memory writes=0
```

4. Compare the results from step 3 and 4 and document all results.

Answer:

Both pen and paper analysis and simulator results **match**, they correspond with each other.

Part 4: Both Temporal and Spatial Locality

Let's modify the logic from part :

sum=0

for (j=0; j < num_ iterations; j++)

for (i = 0; i < num_ elements; i++)

sum=sum+A[i]+B[i]

1. Assume that A[0] is placed in address 0x0 and B[0] is in address 0x100

2. (Paper and Pen Analysis): Experiment with num_elements=8, a four-way set associative cache of size 256-B and line size of 16-B. Vary number of iterations as 1, 5, 10, and 100. For each configuration, obtain the hit rate and classify the misses, record your observations.

Answer:

Size of int = 4 bytes

Line size = 16 bytes

Length of arrays in this case = num_ elements = 8

Block size = 16 bytes
 Cache associativity = 4
 Cache size= 256 bytes
 Number of blocks = Cache lines = $256 / 16 = 16$ blocks
 Number of sets = $16 / 4 = 4$

Assuming, Address bit length= 24 bit
 Bits for [tag , index , offset] = [18 , 2 , 4] = 24
 Extra space for tag storage = 36 bytes

Number of iterations = 1 , 5 , 10 and 100

Variable	Hex Address [tag, index, offset]			Decimal
	tag	index	offset	
	18	2	4	
A[0]	000000 0000 0000 0000	00	0000	0
A[1]	000000 0000 0000 0000	00	0100	4
A[2]	000000 0000 0000 0000	00	1000	8
A[3]	000000 0000 0000 0000	00	1100	12
A[4]	000000 0000 0000 0000	01	0000	16
A[5]	000000 0000 0000 0000	01	0100	20
A[6]	000000 0000 0000 0000	01	1000	24
A[7]	000000 0000 0000 0000	01	1100	28

Variable	Hex Address [tag, index, offset]			Decimal
	tag	index	offset	
	18	2	4	
B[0]	000000 0000 0000 0100	00	0000	256
B[1]	000000 0000 0000 0100	00	0100	260
B[2]	000000 0000 0000 0100	00	1000	264
B[3]	000000 0000 0000 0100	00	1100	268
B[4]	000000 0000 0000 0100	01	0000	272
B[5]	000000 0000 0000 0100	01	0100	276
B[6]	000000 0000 0000 0100	01	1000	280
B[7]	000000 0000 0000 0100	01	1100	284

Case 1: Number of iterations = 1
 Number of cache lines = Cache size / Cache line size = $256B / 16B = 16$ cache lines
 No. of sets = $16 / 4 = 4$

All misses here are Compulsory Misses

Variable	Operation: [Variable Address]	Hit/Miss		Variable	Operation: [Variable Address]	Hit/Miss
A[0]	R: 0x0	Miss		B[0]	R: 0x100	Miss
A[1]	R: 0x4	Hit		B[1]	R: 0x104	Hit
A[2]	R: 0x8	Hit		B[2]	R: 0x108	Hit
A[3]	R: 0xC	Hit		B[3]	R: 0x10C	Hit
A[4]	R: 0x10	Miss		B[4]	R: 0x110	Miss
A[5]	R: 0x14	Hit		B[5]	R: 0x114	Hit
A[6]	R: 0x18	Hit		B[6]	R: 0x118	Hit
A[7]	R: 0x1C	Hit		B[7]	R: 0x11C	Hit

Total traces (read count) = 16

Hits = 12

Misses = 4 (Compulsory Misses)

Memory reads = 4

Hit rate = $12 / 16 = 0.75$

Case 2: Number of iterations = 5

From the second iteration (no. of iterations 2 to 5) all the read operations are a cache hit, because in the first iteration all the memory blocks were brought into cache.

Total traces (read count) = 80

Hits = 76

Misses = 4 (Compulsory Misses)

Memory reads = 4

Hit rate = $76 / 80 = 0.95$

Case 3: Number of iterations = 10

From the second iteration (no. of iterations 2 to 10) all the read operations are a cache hit, because in the first iteration all the memory blocks were brought into cache.

Total traces (read count) = 160

Hits = 156

Misses = 4 (Compulsory Misses)

Memory reads = 4

Hit rate = $156 / 160 = 0.975$

Case 2: Number of iterations = 100

From the second iteration (no. of iterations 2 to 100) all the read operations are a cache hit, because in the first iteration all the memory blocks were brought into cache.

Total traces (read count) = 1600

Hits = 1596

Misses = 4 (Compulsory Misses)

Memory reads = 4

Hit rate = $1596 / 1600 = 0.9975$

3. (Simulation): hand craft the traces for each of the configurations in Step 2 and run this through the simulator you created in Part 2.

Answer:

Here are the traces for the simulation:

R:0

R:4

R:8

R:C

R:10

R:14

R:18

R:1C

R:100

R:104

R:108

R:10C

R:110

R:114

R:118

R:11C

Case 1:

```

neelanjangoswami@Neelanjans-Air Downloads % ./cachesim_final -b 16 -c 256 -w r -a 4 < memory_trace.txt
0KB 4-way associative cache:
Block size = 16 bytes
Number of [sets,blocks] = [4,16]
Extra space for tag storage = 36 bytes (14.06%)
Bits for [tag,index,offset] = [18, 2, 4] = 24
Write policy = Write-back

Hex Binary Address          Set Blk   Memory
Address (tag/index/offset)  Tag Index off Way UWay Read Write
=====
0R 0000000000 00000000 00000 0 0 0 -1 0 1 0
4R 0000000000 00000000 00100 0 0 4 0 -2 1 0
8R 0000000000 00000000 01000 0 0 8 0 -2 1 0
cR 0000000000 00000000 01100 0 0 12 0 -2 1 0
10R 0000000000 00000001 00000 0 1 0 -1 0 1 0
14R 0000000000 00000001 00100 0 1 4 0 -2 1 0
18R 0000000000 00000001 01000 0 1 8 0 -2 1 0
1cR 0000000000 00000001 01100 0 1 12 0 -2 1 0
100R 0000000100 00000000 00000 4 0 0 -1 0 1 0
104R 0000000100 00000000 00100 4 0 4 0 -2 1 0
108R 0000000100 00000000 01000 4 0 8 0 -2 1 0
10cR 0000000100 00000000 01100 4 0 12 0 -2 1 0
110R 0000000100 00000001 00000 4 1 0 -1 0 1 0
114R 0000000100 00000001 00100 4 1 4 0 -2 1 0
118R 0000000100 00000001 01000 4 1 8 0 -2 1 0
11cR 0000000100 00000001 01100 4 1 12 0 -2 1 0
nref=16, nread=16, nwrite=0
hits = 12, hit rate = 0.75
misses = 4, miss rate = 0.25
main memory reads=16, main memory writes=0

```

Case 2:

```

104R 0000000100 00000000 00100 4 0 4 2 -2 1 0
108R 0000000100 00000000 01000 4 0 8 2 -2 1 0
10cR 0000000100 00000000 01100 4 0 12 2 -2 1 0
110R 0000000100 00000001 00000 4 1 0 2 -2 1 0
114R 0000000100 00000001 00100 4 1 4 2 -2 1 0
118R 0000000100 00000001 01000 4 1 8 2 -2 1 0
11cR 0000000100 00000001 01100 4 1 12 2 -2 1 0
0R 0000000000 00000000 00000 0 0 0 1 -2 1 0
4R 0000000000 00000000 00100 0 0 4 1 -2 1 0
8R 0000000000 00000000 01000 0 0 8 1 -2 1 0
cR 0000000000 00000000 01100 0 0 12 1 -2 1 0
10R 0000000000 00000001 00000 0 1 0 1 -2 1 0
14R 0000000000 00000001 00100 0 1 4 1 -2 1 0
18R 0000000000 00000001 01000 0 1 8 1 -2 1 0
1cR 0000000000 00000001 01100 0 1 12 1 -2 1 0
100R 0000000100 00000000 00000 4 0 0 2 -2 1 0
104R 0000000100 00000000 00100 4 0 4 2 -2 1 0
108R 0000000100 00000000 01000 4 0 8 2 -2 1 0
10cR 0000000100 00000000 01100 4 0 12 2 -2 1 0
110R 0000000100 00000001 00000 4 1 0 2 -2 1 0
114R 0000000100 00000001 00100 4 1 4 2 -2 1 0
118R 0000000100 00000001 01000 4 1 8 2 -2 1 0
11cR 0000000100 00000001 01100 4 1 12 2 -2 1 0
0R 0000000000 00000000 00000 0 0 0 1 -2 1 0
4R 0000000000 00000000 00100 0 0 4 1 -2 1 0
8R 0000000000 00000000 01000 0 0 8 1 -2 1 0
cR 0000000000 00000000 01100 0 0 12 1 -2 1 0
10R 0000000000 00000001 00000 0 1 0 1 -2 1 0
14R 0000000000 00000001 00100 0 1 4 1 -2 1 0
18R 0000000000 00000001 01000 0 1 8 1 -2 1 0
1cR 0000000000 00000001 01100 0 1 12 1 -2 1 0
100R 0000000100 00000000 00000 4 0 0 2 -2 1 0
104R 0000000100 00000000 00100 4 0 4 2 -2 1 0
108R 0000000100 00000000 01000 4 0 8 2 -2 1 0
10cR 0000000100 00000000 01100 4 0 12 2 -2 1 0
110R 0000000100 00000001 00000 4 1 0 2 -2 1 0
114R 0000000100 00000001 00100 4 1 4 2 -2 1 0
118R 0000000100 00000001 01000 4 1 8 2 -2 1 0
11cR 0000000100 00000001 01100 4 1 12 2 -2 1 0
nref=80, nread=80, nwrite=0
hits = 76, hit rate = 0.95
misses = 4, miss rate = 0.05
main memory reads=80, main memory writes=0

```

Case 3:


```

104R 00000000100 00000000 00100 4 0 4 2 -2 1 0
108R 00000000100 00000000 01000 4 0 8 2 -2 1 0
10cR 00000000100 00000000 01100 4 0 12 2 -2 1 0
110R 00000000100 00000001 00000 4 1 0 2 -2 1 0
114R 00000000100 00000001 00100 4 1 4 2 -2 1 0
118R 00000000100 00000001 01000 4 1 8 2 -2 1 0
11cR 00000000100 00000001 01100 4 1 12 2 -2 1 0
0R 00000000000 00000000 00000 0 0 0 1 -2 1 0
4R 00000000000 00000000 00100 0 0 4 1 -2 1 0
8R 00000000000 00000000 01000 0 0 8 1 -2 1 0
cR 00000000000 00000000 01100 0 0 12 1 -2 1 0
10R 00000000000 00000001 00000 0 1 0 1 -2 1 0
14R 00000000000 00000001 00100 0 1 4 1 -2 1 0
18R 00000000000 00000001 01000 0 1 8 1 -2 1 0
1cR 00000000000 00000001 01100 0 1 12 1 -2 1 0
100R 00000000100 00000000 00000 4 0 0 2 -2 1 0
104R 00000000100 00000000 00100 4 0 4 2 -2 1 0
108R 00000000100 00000000 01000 4 0 8 2 -2 1 0
10cR 00000000100 00000000 01100 4 0 12 2 -2 1 0
110R 00000000100 00000001 00000 4 1 0 2 -2 1 0
114R 00000000100 00000001 00100 4 1 4 2 -2 1 0
118R 00000000100 00000001 01000 4 1 8 2 -2 1 0
11cR 00000000100 00000001 01100 4 1 12 2 -2 1 0
0R 00000000000 00000000 00000 0 0 0 1 -2 1 0
4R 00000000000 00000000 00100 0 0 4 1 -2 1 0
8R 00000000000 00000000 01000 0 0 8 1 -2 1 0
cR 00000000000 00000000 01100 0 0 12 1 -2 1 0
10R 00000000000 00000001 00000 0 1 0 1 -2 1 0
14R 00000000000 00000001 00100 0 1 4 1 -2 1 0
18R 00000000000 00000001 01000 0 1 8 1 -2 1 0
1cR 00000000000 00000001 01100 0 1 12 1 -2 1 0
100R 00000000100 00000000 00000 4 0 0 2 -2 1 0
104R 00000000100 00000000 00100 4 0 4 2 -2 1 0
108R 00000000100 00000000 01000 4 0 8 2 -2 1 0
10cR 00000000100 00000000 01100 4 0 12 2 -2 1 0
110R 00000000100 00000001 00000 4 1 0 2 -2 1 0
114R 00000000100 00000001 00100 4 1 4 2 -2 1 0
118R 00000000100 00000001 01000 4 1 8 2 -2 1 0
11cR 00000000100 00000001 01100 4 1 12 2 -2 1 0

```

nref=160, nread=160, nwrite=0

hits = 156, hit rate = 0.97

misses = 4, miss rate = 0.03

main memory reads=160, main memory writes=0

neelaniangoswami@Neelaniangoswami-Air: Downloads %

Case 4:

```

104R 0000000100 00000000 00100 4 0 4 2 -2 1 0
108R 0000000100 00000000 01000 4 0 8 2 -2 1 0
10cR 0000000100 00000000 01100 4 0 12 2 -2 1 0
110R 0000000100 00000001 00000 4 1 0 2 -2 1 0
114R 0000000100 00000001 00100 4 1 4 2 -2 1 0
118R 0000000100 00000001 01000 4 1 8 2 -2 1 0
11cR 0000000100 00000001 01100 4 1 12 2 -2 1 0
0R 0000000000 00000000 00000 0 0 0 1 -2 1 0
4R 0000000000 00000000 00100 0 0 4 1 -2 1 0
8R 0000000000 00000000 01000 0 0 8 1 -2 1 0
cR 0000000000 00000000 01100 0 0 12 1 -2 1 0
10R 0000000000 00000001 00000 0 1 0 1 -2 1 0
14R 0000000000 00000001 00100 0 1 4 1 -2 1 0
18R 0000000000 00000001 01000 0 1 8 1 -2 1 0
1cR 0000000000 00000001 01100 0 1 12 1 -2 1 0
100R 0000000100 00000000 00000 4 0 0 2 -2 1 0
104R 0000000100 00000000 00100 4 0 4 2 -2 1 0
108R 0000000100 00000000 01000 4 0 8 2 -2 1 0
10cR 0000000100 00000000 01100 4 0 12 2 -2 1 0
110R 0000000100 00000001 00000 4 1 0 2 -2 1 0
114R 0000000100 00000001 00100 4 1 4 2 -2 1 0
118R 0000000100 00000001 01000 4 1 8 2 -2 1 0
11cR 0000000100 00000001 01100 4 1 12 2 -2 1 0
0R 0000000000 00000000 00000 0 0 0 1 -2 1 0
4R 0000000000 00000000 00100 0 0 4 1 -2 1 0
8R 0000000000 00000000 01000 0 0 8 1 -2 1 0
cR 0000000000 00000000 01100 0 0 12 1 -2 1 0
10R 0000000000 00000001 00000 0 1 0 1 -2 1 0
14R 0000000000 00000001 00100 0 1 4 1 -2 1 0
18R 0000000000 00000001 01000 0 1 8 1 -2 1 0
1cR 0000000000 00000001 01100 0 1 12 1 -2 1 0
100R 0000000100 00000000 00000 4 0 0 2 -2 1 0
104R 0000000100 00000000 00100 4 0 4 2 -2 1 0
108R 0000000100 00000000 01000 4 0 8 2 -2 1 0
10cR 0000000100 00000000 01100 4 0 12 2 -2 1 0
110R 0000000100 00000001 00000 4 1 0 2 -2 1 0
114R 0000000100 00000001 00100 4 1 4 2 -2 1 0
118R 0000000100 00000001 01000 4 1 8 2 -2 1 0
11cR 0000000100 00000001 01100 4 1 12 2 -2 1 0
nref=1600, nread=1600, nwrite=0
hits = 1596, hit rate = 1.00
misses = 4, miss rate = 0.00
main memory reads=1600, main memory writes=0

```

4. Compare the results from step 3 and 4 and document all results.

Answer:

Both pen and paper analysis and simulator results **match**, they correspond with each other.

5. What you expect to change if sum and both arrays are uint_8 instead of int?

Answer:

When you change an integer from an int to a uint_8, the size of the integer shrinks from 4 bytes to 1 byte. This means that the combined memory used by both arrays reduces from 32 bytes to 8 bytes, and these 8 bytes are allocated consecutively.

As the size of the int decreases, the entire array is loaded into the cache when it's first accessed. Consequently, the number of cache misses drops to just 2 compulsory misses. After that, each read request becomes a cache hit.

Part 5: Both Temporal and Spatial Locality with Strides

For the following logic:

```
sum=0 // int  
  
for (j=0; J < num_ iterations; j++)  
  
for(i = 0; i < num_ elements; i=i+stride)  
  
sum=sum+A[i]
```

1. Assume that A[0] is placed in address 0x0

Use 'stride.cpp' for simulation for all of the following problems.

2. Experiment (both paper and pen and in simulation) with num_elements=128, a four-way set associative cache of size 256-B and line size of 16-B and number of iterations of 100. Vary the stride size as 1,2,4,8,16,21,64.

a. Obtain the hit rate and classify the misses

Answer:

Size of int = 4 bytes

Line size = 16 bytes

Length of arrays in this case = num_elements = 8

Block size = 16 bytes

Cache associativity = 4

Cache size= 256 bytes

Number of blocks = Cache lines = 256 / 16 = 16 blocks

Number of sets = 16 / 4 = 4

Assuming, Address bit length= 24 bit

Bits for [tag , index , offset] = [18 , 2 , 4] = 24

Extra space for tag storage = 36 bytes

Stride = 1 , 2 , 4 , 8 , 16 , 21 , 64

Variable	Hex Address [tag, index, offset]			Decimal
	tag	index	offset	
	18	2	4	
A[0]	000000 0000 0000 0000	00	0000	0
A[1]	000000 0000 0000 0000	00	0100	4
A[2]	000000 0000 0000 0000	00	1000	8
A[3]	000000 0000 0000 0000	00	1100	12
A[4]	000000 0000 0000 0000	01	0000	16
A[5]	000000 0000 0000 0000	01	0100	20
A[6]	000000 0000 0000 0000	01	1000	24
A[7]	000000 0000 0000 0000	01	1100	28
.
.
.
A[124]	000000 0000 0000 0111	11	0000	496
A[125]	000000 0000 0000 0111	11	0100	500
A[126]	000000 0000 0000 0111	11	1000	504
A[127]	000000 0000 0000 0111	11	1100	508

Case 1: Stride = 1

Value of i = 0 , 1 , 2 , 3 , ... , 125 , 126 , 127

Cache is filled in this process:

	Way 00	Way 01	Way 10	Way 11
00	A[0...3]	A[4...7]	A[8...11]	A[12...15]
01	A[16...19]	A[20...23]	A[24...27]	A[28...31]
10	A[32...35]	A[36...39]	A[40...43]	A[44...47]
11	A[48...51]	A[52...55]	A[56...59]	A[60...63]

Therefore, in 1 loop:

Total Traces (read count) = 128

Misses = 32 misses (compulsory/cold miss)

Hits = 96

Number of iterations = 100

Total Read Count = 12,800

Cache Misses = 3200 :-

Cold Misses = 32

Capacity Misses = 3168

Hits = 9600

Hit rate = $9600/12800 = 0.75$

```
neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 12800
Total Hits: 9600
Total Misses: 3200
Hit Rate: 75%
Miss Rate: 25%
```

Both pen and paper analysis and simulator results **match**, they correspond with each other.

Case 2: Stride = 2

Value of $i = 0, 2, 4, \dots, 124, 126$

Cache is filled in this process:

	Way 00	Way 01	Way 10	Way 11
00	A[0...3]	A[16...19]	A[32...35]	A[48...51]
01	A[4...7]	A[20...23]	A[36...39]	A[52...55]
10	A[8...11]	A[24...27]	A[40...43]	A[56...59]
11	A[12...15]	A[28...31]	A[44...47]	A[60...63]

Therefore, in 1 loop:

Total Traces (read count) = 64

Misses = 32 misses (compulsory/cold miss)

Hits = 32

Number of iterations = 100

Total Read Count = 6400

Cache Misses = 3200 :-

Cold Misses = 32

Capacity Misses = 3168

Hits = 3200

Hit rate = $3200/6400 = 0.5$

```
neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 6400
Total Hits: 3200
Total Misses: 3200
Hit Rate: 50%
Miss Rate: 50%
```

Both pen and paper analysis and simulator results **match**, they correspond with each other.

Case 3: Stride = 4

Value of $i = 0, 4, 8, 12, \dots, 120, 124$

Cache is filled in this process:

	Way 00	Way 01	Way 10	Way 11
00	A[0...3]	A[16...19]	A[32...35]	A[48...51]
01	A[4...7]	A[20...23]	A[36...39]	A[52...55]
10	A[8...11]	A[24...27]	A[40...43]	A[56...59]
11	A[12...15]	A[28...31]	A[44...47]	A[60...63]

Therefore, in 1 loop:

Total Traces (read count) = 32

Misses = 32 misses (compulsory/cold miss)

Hits = 0

Number of iterations = 100

Total Read Count = 3200

Cache Misses = 3200 :-

Cold Misses = 32

Capacity Misses = 3168

Hits = 0

Hit rate = $0/3200 = 0$

```
neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 3200
Total Hits: 0
Total Misses: 3200
Hit Rate: 0%
Miss Rate: 100%
```

Both pen and paper analysis and simulator results **match, they correspond with each other.**

Case 4: Stride = 8

Value of i = 0 , 8 , 16 , ... , 116 , 120

Cache is filled in this process:

	Way 00	Way 01	Way 10	Way 11
00	A[0...3]	A[16...19]	A[32...35]	A[48...51]
01				
10	A[8...11]	A[24...27]	A[40...43]	A[56...59]
11				

Therefore, in 1 loop:

Total Traces (read count) = 16

Misses = 16 misses (compulsory/cold miss)

Hits = 0

Number of iterations = 100

Total Read Count = 1600

Cache Misses = 1600 :-

Cold Misses = 16

Capacity Misses = 1584

Hits = 0

Hit rate = $0/1600 = 0$

```
neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 1600
Total Hits: 0
Total Misses: 1600
Hit Rate: 0%
Miss Rate: 100%
```

Both pen and paper analysis and simulator results **match, they correspond with each other.**

Case 5: Stride = 16

Value of $i = 0, 16, 32, \dots, 96, 112$

Cache is filled in this process:

	Way 00	Way 01	Way 10	Way 11
00	A[0...3]	A[16...19]	A[32...35]	A[48...51]
01				
10				
11				

Therefore, in 1 loop:

Total Traces (read count) = 8

Misses = 8 misses (compulsory/cold miss)

Hits = 0

Number of iterations = 100

Total Read Count = 800

Cache Misses = 800 :-

Cold Misses = 8

Capacity Misses = 792

Hits = 0

Hit rate = $0/800 = 0$

```

neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 800
Total Hits: 0
Total Misses: 800
Hit Rate: 0%
Miss Rate: 100%

```

Both pen and paper analysis and simulator results **match, they correspond with each other.**

Case 6: Stride = 21

Value of i = 0 , 21 , 42 , 63 , 84 , 105 , 126

Cache is filled in this process:

	Way 00	Way 01	Way 10	Way 11
00	A[0...3]			
01	A[20...23]	A[84...87]		
10	A[40...43]	A[104...107]		
11	A[60...63]	A[124...127]		

Therefore, in 1 loop:

Total Traces (read count) = 7

Misses = 7 misses (compulsory/cold miss)

Hits = 0

Number of iterations = 100

Total Read Count = 700

Cache Misses = 99 :-

Cold Misses = 99

Hits = 601

Hit rate = $693/700 = 0.14$

```

neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 700
Total Hits: 99
Total Misses: 601
Hit Rate: 14.1429%
Miss Rate: 85.8571%

```

Both pen and paper analysis and simulator results **match, they correspond with each other.**

Case 7: Stride = 64

Value of i = 0 , 64

Cache is filled in this process:

	Way 00	Way 01	Way 10	Way 11
--	--------	--------	--------	--------

00	A[0...3]	A[64...67]		
01				
10				
11				

Therefore, in 1 loop:

Total Traces (read count) = 2

Misses = 2 misses (compulsory/cold miss)

Hits = 0

Number of iterations = 100

Total Read Count = 200

Cache Misses = 200 :-

Cold Misses = 200

Hits = 0

Hit rate = $0/200 = 0$

```
neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 200
Total Hits: 0
Total Misses: 200
Hit Rate: 0%
Miss Rate: 100%
```

Both pen and paper analysis and simulator results **match, they correspond with each other.**

3. Repeat Step 2 for number of iterations=10 and 50

Answer:

Initial steps same as the previous answer.

Case 1: Stride = 1

Value of i = 0 , 1 , 2 , 3 , ... , 125 , 126 , 127

Cache is filled in this process:

	Way 00	Way 01	Way 10	Way 11
00	A[0...3]	A[4...7]	A[8...11]	A[12...15]
01	A[16...19]	A[20...23]	A[24...27]	A[28...31]
10	A[32...35]	A[36...39]	A[40...43]	A[44...47]
11	A[48...51]	A[52...55]	A[56...59]	A[60...63]

Therefore, in 1 loop:

Total Traces (read count) = 128
Misses = 32 misses (compulsory/cold miss)
Hits = 96

Number of iterations = 10
Total Read Count = 1280
Cache Misses = 320 :-
 Cold Misses = 32
 Capacity Misses = 288
Hits = 960
Hit rate = $960/1280 = 0.75$

```
neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 1280
Total Hits: 960
Total Misses: 320
Hit Rate: 75%
Miss Rate: 25%
```

Number of iterations = 50
Total Read Count = 6400
Cache Misses = 1600 :-
 Cold Misses = 32
 Capacity Misses = 1568
Hits = 4800
Hit rate = $4800/6400 = 0.75$

```
neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 6400
Total Hits: 4800
Total Misses: 1600
Hit Rate: 75%
Miss Rate: 25%
```

Both pen and paper analysis and simulator results **match, they correspond with each other.**

Case 2: Stride = 2
Value of i = 0 , 2 , 4 , ... , 124 , 126

Cache is filled in this process:

	Way 00	Way 01	Way 10	Way 11
00	A[0...3]	A[16...19]	A[32...35]	A[48...51]
01	A[4...7]	A[20...23]	A[36...39]	A[52...55]
10	A[8...11]	A[24...27]	A[40...43]	A[56...59]
11	A[12...15]	A[28...31]	A[44...47]	A[60...63]

Therefore, in 1 loop:

Total Traces (read count) = 64
Misses = 32 misses (compulsory/cold miss)
Hits = 32

Number of iterations = 10

Total Read Count = 640

Cache Misses = 320 :-

Cold Misses = 32

Capacity Misses = 288

Hits = 320

Hit rate = $320/640 = 0.5$

```
neelanjangoswami@neelanjans-Air Downloads % ./stride
Total count: 640
Total Hits: 320
Total Misses: 320
Hit Rate: 50%
Miss Rate: 50%
```

Number of iterations = 50

Total Read Count = 3200

Cache Misses = 1600 :-

Cold Misses = 32

Capacity Misses = 1568

Hits = 1600

Hit rate = $1600/3200 = 0.5$

```
neelanjangoswami@neelanjans-Air Downloads % ./stride
Total count: 3200
Total Hits: 1600
Total Misses: 1600
Hit Rate: 50%
Miss Rate: 50%
```

Both pen and paper analysis and simulator results **match, they correspond with each other.**

Case 3: Stride = 4

Value of i = 0 , 4 , 8 , 12 , ... , 120 , 124

Cache is filled in this process:

	Way 00	Way 01	Way 10	Way 11
00	A[0...3]	A[16...19]	A[32...35]	A[48...51]
01	A[4...7]	A[20...23]	A[36...39]	A[52...55]
10	A[8...11]	A[24...27]	A[40...43]	A[56...59]
11	A[12...15]	A[28...31]	A[44...47]	A[60...63]

Therefore, in 1 loop:

Total Traces (read count) = 32
Misses = 32 misses (compulsory/cold miss)
Hits = 0

Number of iterations = 10
Total Read Count = 320
Cache Misses = 320 :-
 Cold Misses = 32
 Capacity Misses = 288
Hits = 0
Hit rate = $0/320 = 0$

```
neelanjanganoswami@Neelanjans-Air Downloads % ./stride
Total count: 320
Total Hits: 0
Total Misses: 320
Hit Rate: 0%
Miss Rate: 100%
```

Number of iterations = 50
Total Read Count = 1600
Cache Misses = 1600 :-
 Cold Misses = 32
 Capacity Misses = 1568
Hits = 0
Hit rate = $0/1600 = 0$

```
neelanjanganoswami@Neelanjans-Air Downloads % ./stride
Total count: 1600
Total Hits: 0
Total Misses: 1600
Hit Rate: 0%
Miss Rate: 100%
```

Both pen and paper analysis and simulator results **match, they correspond with each other.**

Case 4: Stride = 8
Value of $i = 0, 8, 16, \dots, 116, 120$

Cache is filed in this process:

	Way 00	Way 01	Way 10	Way 11
00	A[0...3]	A[16...19]	A[32...35]	A[48...51]
01				
10	A[8...11]	A[24...27]	A[40...43]	A[56...59]
11				

Therefore, in 1 loop:

Total Traces (read count) = 16
Misses = 16 misses (compulsory/cold miss)
Hits = 0

Number of iterations = 10
Total Read Count = 160
Cache Misses = 160 :-
 Cold Misses = 16
 Capacity Misses = 144
Hits = 0
Hit rate = $0/160 = 0$

```
neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 160
Total Hits: 0
Total Misses: 160
Hit Rate: 0%
Miss Rate: 100%
```

Number of iterations = 50
Total Read Count = 800
Cache Misses = 800 :-
 Cold Misses = 16
 Capacity Misses = 784
Hits = 0
Hit rate = $0/800 = 0$

```
neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 800
Total Hits: 0
Total Misses: 800
Hit Rate: 0%
Miss Rate: 100%
```

Both pen and paper analysis and simulator results **match, they correspond with each other.**

Case 5: Stride = 16
Value of i = 0 , 16 , 32 , ... , 96 , 112

Cache is filled in this process:

	Way 00	Way 01	Way 10	Way 11
00	A[0...3]	A[16...19]	A[32...35]	A[48...51]
01				
10				
11				

Therefore, in 1 loop:

Total Traces (read count) = 8
Misses = 8 misses (compulsory/cold miss)
Hits = 0

Number of iterations = 10

Total Read Count = 80

Cache Misses = 80 :-

Cold Misses = 8

Capacity Misses = 72

Hits = 0

Hit rate = $0/80 = 0$

```
neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 80
Total Hits: 0
Total Misses: 80
Hit Rate: 0%
Miss Rate: 100%
```

Number of iterations = 50

Total Read Count = 400

Cache Misses = 400 :-

Cold Misses = 8

Capacity Misses = 392

Hits = 0

Hit rate = $0/400 = 0$

```
neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 400
Total Hits: 0
Total Misses: 400
Hit Rate: 0%
Miss Rate: 100%
```

Both pen and paper analysis and simulator results **match**, they correspond with each other.

Case 6: Stride = 21

Value of i = 0 , 21 , 42 , 63 , 84 , 105 , 126

Cache is filled in this process:

	Way 00	Way 01	Way 10	Way 11
00	A[0...3]			
01	A[20...23]	A[84...87]		
10	A[40...43]	A[104...107]		
11	A[60...63]	A[124...127]		

Therefore, in 1 loop:

Total Traces (read count) = 7
Misses = 7 misses (compulsory/cold miss)
Hits = 0

Number of iterations = 10

Total Read Count = 70

Cache Misses = 7 :-

Cold Misses = 7

Hits = 63

Hit rate = $63/70 = 0.9$

```
neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 70
Total Hits: 9
Total Misses: 61
Hit Rate: 12.8571%
Miss Rate: 87.1429%
```

Number of iterations = 50

Total Read Count = 350

Cache Misses = 7 :-

Cold Misses = 7

Hits = 343

Hit rate = $343/350 = 0.98$

```
neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 350
Total Hits: 49
Total Misses: 301
Hit Rate: 14%
Miss Rate: 86%
```

Both pen and paper analysis and simulator results **match**, they correspond with each other.

Case 7: Stride = 64

Value of $i = 0, 64$

Cache is filled in this process:

	Way 00	Way 01	Way 10	Way 11
00	A[0...3]	A[64...67]		
01				
10				
11				

Therefore, in 1 loop:

Total Traces (read count) = 2

Misses = 2 misses (compulsory/cold miss)

Hits = 0

Number of iterations = 10

Total Read Count = 20

Cache Misses = 20 :-

Cold Misses = 20

Hits = 0

Hit rate = $0/20 = 0$

```
neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 20
Total Hits: 0
Total Misses: 20
Hit Rate: 0%
Miss Rate: 100%
```

Number of iterations = 50

Total Read Count = 200

Cache Misses = 200 :-

Cold Misses = 2

Hits = 0

Hit rate = $0/200 = 0$

```
neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 100
Total Hits: 0
Total Misses: 100
Hit Rate: 0%
Miss Rate: 100%
```

Both pen and paper analysis and simulator results **match, they correspond with each other.**

4. How does varying number of iterations and stride value affect the hit rate? How does this relate to spatial and temporal locality?

Answer:

When you increase the stride value, initially it negatively impacts the hit rate, causing it to decrease. Once it becomes equal to the size of an integer, the hit rate drops to zero. When you continue to increase the stride in multiples of 2, the hit rate remains at zero because all the accessed elements are mapped to the same cache set. On the other hand, if you increase the stride so much that the cache isn't fully filled, after the first iteration, it's always a hit because there's no conflict. As you increase the number of iterations, hits generally increase because there are no cold misses. However, this outcome depends significantly on the stride value. In cases of spatial locality, a smaller stride value increases the cache hit rate, and vice versa. This happens because the entire block is prefetched, and adjacent memory locations are accessed quickly.

In terms of temporal locality, the cache status is crucial. I observed that when the stride value is chosen in a way that the memory addresses are mapped to the same cache set, there are very few

cache hits. This is because there's a conflict miss every time. The best scenario occurs when the stride is either too large or not a power of 2.

Part 6: Cache Architecture and Replacement Policies

1. For the same logic in Part 5, experiment (both paper and pen and in simulation) with num_elements=128, a cache size of size 256-B and line size of 16-B, a stride of 1 and number of iterations of 100. Obtain the hit rate and identify the final cache content for the following cache architectures

a. Direct Mapped

Answer:

Length of array = num_elements = 128

Cache size = 256 bytes

Block size = 16 bytes

Stride = 1

No. of blocks = $256 / 16 = 16$

Num_iterations = 100

Associativity = 1

No. of sets = 16

[Tag , Index , Offset] = [16 , 4 , 4] = 24 bits

After number of iterations = 100, cache state:

Index	Value
0000	A[64, 65, 66, 67]
0001	A[68, 69, 70, 71]
.	.
.	.
.	.
1110	A[120, 121, 122, 123]
1111	A[124, 125, 126, 127]

In one loop:

Total traces (read count) = 128

Cache misses = 32 (compulsory/cold miss)

Hits = 96

Therefore,

For number of iterations = 100:

Total count = 12800

Cache misses = 3200

Cold misses = 32

Hit = 9600

Hit rate = $9600/12800 = 0.75$

```
neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 12800
Total Hits: 9600
Total Misses: 3200
Hit Rate: 75%
Miss Rate: 25%
```

Both pen and paper analysis and simulator results **match, they correspond with each other.**

b. 2-way Set Associative

Answer:

Initial steps same as previous answer.

Associativity = 2

No. of sets = $16 / 2 = 8$

[Tag , Index , Offset] = [17 , 3 , 4] = 24 bits

After number of iterations = 100, cache state:

Index	Value (Way 0)	Value (Way 1)
000	A[64, 65, 66, 67]	A[96, 97, 98, 99]
001	A[68, 69, 70, 71]	A[100, 101, 102, 103]
.	.	.
.	.	.
.	.	.
110	A[88, 89, 90, 91]	A[120, 121, 122, 123]
111	A[92, 93, 94, 95]	A[124, 125, 126, 127]

In one loop:

Total traces (read count) = 128

Cache misses = 32 (compulsory/cold miss)

Hits = 96

Therefore,

For number of iterations = 100:

Total count = 12800

Cache misses = 3200

Cold misses = 32

Capacity misses = 3168
Hit = 9600
Hit rate = 9600/12800 = 0.75

Both pen and paper analysis and simulator results **match, they correspond with each other.**

c. 4-way Set Associative

Answer:

Initial steps same as previous answer.

Associativity = 4
No. of sets = $16 / 2 = 4$
[Tag , Index , Offset] = [18 , 2 , 4] = 24 bits

After number of iterations = 100, cache state:

Index	Value (Way 00)	Value (Way 01)	Value (Way 10)	Value (Way 11)
00	A[64, 65, 66, 67]	A[80, 81, 82, 83]	A[96, 97, 98, 99]	A[112, 113, 114, 115]
01	A[68, 69, 70, 71]	A[84, 85, 86, 87]	A[100, 101, 102, 103]	A[116, 117, 118, 119]
10	A[72, 73, 74, 75]	A[88, 89, 90, 91]	A[104, 105, 106, 107]	A[120, 121, 122, 123]
11	A[76, 77, 78, 79]	A[92, 93, 94, 95]	A[108, 109, 110, 111]	A[124, 125, 126, 127]

In one loop:
Total traces (read count) = 128
Cache misses = 32 (compulsory/cold miss)
Hits = 96

Therefore,
For number of iterations = 100:
Total count = 12800
Cache misses = 3200
Cold misses = 32
Capacity misses = 3168
Hit = 9600
Hit rate = 9600/12800 = 0.75

```
neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 12800
Total Hits: 9600
Total Misses: 3200
Hit Rate: 75%
Miss Rate: 25%
```

Both pen and paper analysis and simulator results **match, they correspond with each other.**

d. Fully-Associative Cache

Answer:

Number of blocks = 16

[Tag , Offset] = [20 , 4] = 24 bits

After number of iterations = 100, cache state:

Block	Value
1	A[64, 65, 66, 67]
2	A[68, 69, 70, 71]
.	.
.	.
.	.
15	A[120, 121, 122, 123]
16	A[124, 125, 126, 127]

In one loop:

Total traces (read count) = 128

Cache misses = 32 (Compulsory / Cold misses)

Hits = 96

Therefore,

For number of iterations = 100:

Total count = 12800

Cache misses = 3200

Cold misses = 32

Capacity misses = 3168

Hit = 9600

Hit rate = 9600/12800 = 0.75

```
neelanjangoswami@Neelanjans-Air Downloads % ./stride
Total count: 12800
Total Hits: 9600
Total Misses: 3200
Hit Rate: 75%
Miss Rate: 25%
```

Both pen and paper analysis and simulator results **match, they correspond with each other.**