

ASSIGNMENT 3. BINARY CLASSIFICATION WITH LOGISTIC REGRESSION AND k -NEAREST NEIGHBOURS

Due Date: February 25, 11:30 pm
Assessment: 8% of the total course mark.

DESCRIPTION:

In this assignment you are required to experiment with logistic regression and k -nearest neighbours for binary classification. You will use the “Wisconsin breast cancer” data set, which is available in scikit-learn. The two classes are “malignant” (the **positive** class) and “benign” (the **negative** class) (if the classes are labeled differently, please change their labels). The number of features is 30.

After you separate the test and training data, you have to standardize the features, (i.e. center and scale them - see slide 16 in Topic 2) in the training data and then apply the transformations to the features in the test data. Then you have to do the following.

- ◇ Train a linear classifier using your implementation of logistic regression. Then plot the ROC (receiver operating characteristics) curve (see page 12 in Topic 6). For this you have to consider all possible thresholds for your classifier based on your test data and compute the sensitivity and specificity in each case. Identify on the curve the point corresponding to the classifier that minimizes the misclassification error.
- ◇ Do the same as above using the implementation of logistic regression provided in scikit-learn.
- ◇ Implement the k -nearest neighbour classifier for each $k = 1, 3, 5, 7, 9$, and select the best k using K -fold cross-validation (with your implementation). Consider the 0/1 loss. Then use the best classifier to compute the test error (i.e., the misclassification rate on the test data set).
- ◇ Do the same as above using the implementation provided by scikit-learn.

You have to write a report to present your results and their discussion. The report must include a table with the test errors for all four models. Use the 0/1 error function. For logistic regression use the classifier that minimizes the 0/1 loss. For the nearest neighbours predictors you have to also specify the cross-validation error for each k , in a separate table. You have to specify the value of K that you choose.

Note: In k - nearest neighbours there are situations where ties might occur. For instance, when two or more training examples that are at the same distance from the current example. In such cases, you must decide how to deal with the ties and you have to document this in your report.

For the two logistic regression models you have to specify the vectors of parameters and to include a plot with the two ROC curves.

How do the results obtained with your implementations compare with the results obtained using scikit-learn? Which of the four models is the best overall? To decide this,

use first the misclassification rate as the metric, next use the F1 score as the metric. Is the winner the same in the two cases?

Besides the report, you have to submit your numpy code. The code has to be modular. Write a function for each of the main tasks. Also, write a function for each task that is executed multiple times (e.g, to compute the average error, to compute the K -fold cross-validation error for a model, etc). The code should include instructive comments.

You may use the following code for feature normalization:

```
import numpy as np
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

Note that the normalization is performed based on the training data. Then the same transformations are applied to the test data in order to be able to use the trained predictor on the test data.

SUBMISSION INSTRUCTIONS:

- Submit the report in pdf format, the python file (with extension “.py”) containing your code, and a short demo video. The video should be 2 min or less. In the video, you should scroll down your code, show that it runs and that it outputs the results for each part of the assignment. Submit the files in the Assignments Box on Avenue.

Instructions to obtain the ROC curve for a logistic regression model:

The classifier that minimizes the misclassification rate is

$$\hat{C}(\mathbf{x}) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases},$$

where $z = \mathbf{w}^T \mathbf{x}$. Other classifiers can be obtained by replacing 0 by some other thresholds θ :

$$\hat{C}(\mathbf{x}) = \begin{cases} 1 & \text{if } z \geq \theta \\ 0 & \text{if } z < \theta \end{cases}.$$

To obtain all possible classifiers for the test set, it is sufficient to consider as values for θ , the values $z[0], z[1], \dots, z[M-1]$, obtained on the test examples (where M is the size of the test set). For this, first sort the array \mathbf{z} . Then take in turn $\theta = z[0]$, $\theta = z[1]$, and so on, up to $\theta = z[M-1]$. For each θ compute the **sensitivity and specificity**. Note that the positive class is the “malignant” class.